

# Logistic Regression

# Logistic Regression

## The Idea

- The biggest disadvantage of the Perceptron is that it never converges if the classes are not perfectly linearly separable.
- We will now take a look at another simple yet more powerful algorithm for **linear and binary classification** problems: logistic regression.
  - Note that, in spite of its name, logistic regression is a model for classification, not regression.
- The Logistic regression is a linear model for binary classification that **can be extended to multiclass classification** via the OvR technique.

# Logistic Regression

## The idea

- Logistic regression is a binary classifier: the output variable  $Y$  has two possible values 0 and 1
- Logistic regression is a probabilistic model: its goal is to model the probability of the positive class (i.e., the class that we want to predict), typically class 1.
  - The term positive event does not necessarily mean good, but refers to the event that we want to predict, for example, the probability that a patient has a certain disease; we can think of the positive event as class label  $y = 1$ .
- Consider a single input observation  $\mathbf{x}$ , which we will represent by a vector of features  $[x_1, x_2, \dots, x_n]$ , we want to know the probability that this observation  $\mathbf{x}$  belongs to the positive class 1,  $P(Y = 1|\mathbf{x})$ .

# Logistic Regression

## Mathematical details

- To explain the idea behind logistic regression as a probabilistic model, let's first introduce the **odds ratio**, which is the odds in favor of a particular event.
- The odds ratio can be written as  $\frac{p}{(1-p)}$ , where  $p$  stands for the probability of the positive event.
- We can then further define the **logit function**, which is simply the logarithm of the odds ratio (log-odds):

$$\text{logit}(p) = \log \frac{p}{(1-p)}$$

- The aim of the logistic regression algorithm is to compute:

$$\text{logit}(p(y = 1|\mathbf{x}))$$


# Logistic Regression

## Probabilistic Models

- The logit function takes input values in  $[0, 1]$  and transforms them to values over the entire real number range, which we can use to express a linear relationship between feature values and the log-odds:

$$\text{logit}(p(y = 1|\mathbf{x})) = w_0x_0 + w_1x_1 + \dots + w_mx_m = \sum_{i=0}^m w_ix_i = \mathbf{w}^T \mathbf{x}$$

where  $p(y = 1|\mathbf{x})$  is the conditional probability that a particular sample belongs to class 1 given its features  $\mathbf{x}$ .



# Logistic Regression

## Probabilistic Models

- What we are actually interested in is predicting the probability that a certain sample belongs to a particular class, which is the inverse form of the logit function: the logistic function, sometimes simply abbreviated as **sigmoid function** due to its characteristic S-shape.


$$\phi(z) = \frac{1}{1 + e^{-z}}$$

where  $z$  is the net input, that is, the linear combination of weights and sample features and can be calculated as

$$z = w_0x_0 + w_1x_1 + \dots + w_mx_m = \mathbf{w}^T \mathbf{x}$$

# Logistic Regression

## Probabilistic Models



- Sigmoid function takes real number values as input and transforms them to values in the range  $[0, 1]$  with an intercept at  $\phi(z) = 0.5$ .

# Logistic Regression

## Threshold function

- The predicted probability can then simply be converted into a binary outcome via a quantizer (unit step function):

$$\hat{y} = \begin{cases} 1 & \text{if } \phi(z) \geq 0.5 \\ 0 & \text{otherwise} \end{cases}$$


- If we look at the preceding sigmoid plot, this is equivalent to the following:

$$\hat{y} = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$



# Logistic Regression

## Workflow



- In Adaline, we used the identity function  $\phi(z) = z$  as the activation function. In logistic regression, this activation function simply becomes the sigmoid function that we defined earlier.

# Logistic Regression

## Utility function

- Let's consider a single observation  $x$ , we can define the output of the sigmoid function as the probability of particular sample to belong to the positive class (i.e, class 1):

$$\phi(z) = p(y = 1|x) = \hat{y} \text{ and } 1 - \phi(z) = p(y = 0|x) = 1 - \hat{y}$$

- We can express the probability  $p(y|x)$  that our classifier produces for one observation as the following:

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

- Let's consider the function  $L$  that we want to maximise when we build a logistic regression model as:

$$L(\mathbf{w}) = P(\mathbf{y}|\mathbf{x}; \mathbf{w})$$

where  $P(\mathbf{y}|\mathbf{x}; \mathbf{w})$  is the probability that our classifier produces given its features  $\mathbf{x}$  parameterized by the weights  $\mathbf{w}$ .

# Logistic Regression

## Utility Function

- Mathematically,

$$L(\mathbf{w}) = P(\mathbf{y}|\mathbf{x};\mathbf{w}) = \prod_{i=1}^n P(y^{(i)}|x^{(i)};\mathbf{w}) = \prod_{i=1}^n (\phi(z^{(i)})^{y^{(i)}} (1-\phi(z^{(i)}))^{1-y^{(i)}})$$

- We perform the  $\log L(\mathbf{w})$ , since, it is easier to maximize the (natural) log of this equation, which is called the log-likelihood function:

$$l(\mathbf{w}) = \log L(\mathbf{w}) = \sum_{i=1}^n \left[ y^{(i)} \log(\phi(z^{(i)})) + (1-y^{(i)}) \log(1 - \phi(z^{(i)})) \right]$$



$$\log_a(b \cdot c) = \log_a(b) + \log_a(c)$$

$$\log_a(b^c) = c \log_a(b)$$

# Logistic Regression

## Cost Function

- We could use an optimization algorithm such as gradient ascent to maximize the log-likelihood function. Alternatively, let's rewrite the log-likelihood as a cost function,  $J$ , that can be minimized using **gradient descent**:

$$J(\mathbf{w}) = \sum_{i=1}^n \left[ -y^{(i)} \log(\phi(z^{(i)})) - (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right]$$

- To get a better grasp on this cost function, let's take a look at the cost that we calculate for one single-sample instance:


$$J(\phi(z); y; \mathbf{w}) = -y \log(\phi(z)) - (1 - y) \log(1 - \phi(z))$$

- Looking at the preceding equation, we can see that:

$$J(\phi(z); y; \mathbf{w}) = \begin{cases} -\log(\phi(z)) & \text{if } y = 1 \\ -\log(1 - \phi(z)) & \text{if } y = 0 \end{cases}$$

# Logistic Regression

## Observations on Cost Function



- The cost approaches 0 (plain blue line) if we correctly predict that a sample belongs to class 1. Similarly, we can see on the  $y$  axis that the cost also approaches 0 if we correctly predict  $y = 0$  (dashed line). However, if the prediction is wrong, the cost goes towards infinity: we penalize wrong predictions with an increasingly larger cost.

# Logistic Regression

## Converting an Adaline Implementation into an Algorithm for Logistic Regression

- If we were to implement logistic regression ourselves, we could simply substitute the cost function  $J$  in our Adaline implementation with the new cost function:

$$J(\mathbf{w}) = \sum_{i=1}^n \left[ -y^{(i)} \log(\phi(z^{(i)})) - (1 - y^{(i)}) \log(1 - \phi(z^{(i)})) \right]$$

- Also, we need to swap the linear activation function with the sigmoid activation and change the threshold function to return class labels 0 and 1 instead of -1 and 1.

# Logistic Regression

## Converting an Adaline Implementation into an Algorithm for Logistic Regression

---

```
class LogisticRegressionGD(object):

    """Logistic Regression Classifier using gradient descent.
    Parameters
    -----
    eta : (float) Learning rate (between 0.0 and 1.0)
    n_iter : (int) Passes over the training dataset.
    random_state : (int) Random number generator seed for random
        weight initialization.

    Attributes
    -----
    w_ : (1d-array) Weights after fitting.
    cost_ : (list) Sum-of-squares cost function value in each
        epoch.
    """
```

# Logistic Regression

## Converting an Adaline Implementation into an Algorithm for Logistic Regression

```
def __init__(self, eta=0.05, n_iter=100, random_state=1):
    self.eta = eta
    self.n_iter = n_iter
    self.random_state = random_state
def fit(self, X, y):
    """ Fit training data.
    Parameters
    -----
    X : ({array-like}, shape = [n_samples, n_features])
        Training vectors, where n_samples is the number of
        samples and n_features is the number of features.
    y : array-like, shape = [n_samples] Target values.
    Returns
    -----
    self : object
    """
```



# Logistic Regression

## Converting an Adaline Implementation into an Algorithm for Logistic Regression

```
rgen = np.random.RandomState(self.random_state)
self.w_ = rgen.normal(loc=0.0, scale=0.01, size=1 +
    X.shape[1])
self.cost_ = []
for i in range(self.n_iter):
    net_input = self.net_input(X)
    output = self.activation(net_input)
    errors = (y - output)
    self.w_[1:] += self.eta * X.T.dot(errors)
    self.w_[0] += self.eta * errors.sum()
    # note that we compute the logistic 'cost' now
    # instead of the sum of squared errors cost
    cost = (-y.dot(np.log(output)) - ((1 -
        y).dot(np.log(1 - output))))
    self.cost_.append(cost)
return self
```

# Logistic Regression

## Converting an Adaline Implementation into an Algorithm for Logistic Regression

---

```
def net_input(self, X):
    """Calculate net input"""
    return np.dot(X, self.w_[1:]) + self.w_[0]

def activation(self, z):
    """Compute logistic sigmoid activation"""
    return 1. / (1. + np.exp(-np.clip(z, -250, 250)))

def predict(self, X):
    """Return class label after unit step"""
    return np.where(self.net_input(X) >= 0.0, 1, 0)
    # equivalent to:
    # return np.where(self.activation(self.net_input(X))
    #                 # >= 0.5, 1, 0)
```

---

# Logistic Regression

## Converting an Adaline Implementation into an Algorithm for Logistic Regression

- When we fit a logistic regression model, we have to keep in mind that it only works for binary classification tasks. So, let us consider only Iris-setosa and Iris-versicolor flowers (classes 0 and 1) and check that our implementation of logistic regression works:


---

```
>>> X_train_01_subset = X_train[(y_train==0)|(y_train==1)]
>>> y_train_01_subset = y_train[(y_train==0)|(y_train==1)]
>>> lrgd = LogisticRegressionGD(eta=0.05, n_iter=1000,
...     random_state=1)
>>> lrgd.fit(X_train_01_subset, y_train_01_subset)
>>> plot_decision_regions(X=X_train_01_subset,
...     y=y_train_01_subset, classifier=lrgd)
>>> plt.xlabel('petal length [standardized]')
>>> plt.ylabel('petal width [standardized]')
>>> plt.legend(loc='upper left')
>>> plt.show()
```

---

# Logistic Regression

Converting an Adaline Implementation into an Algorithm for Logistic Regression



# Logistic Regression

## Training a Logistic Regression Model via SciKit-Learn

- Let's learn how to use scikit-learn's more optimized implementation of logistic regression that also supports multi-class settings off the shelf (OvR by default).


---

```
>>> from sklearn.linear_model import LogisticRegression
>>> lr = LogisticRegression(C=100.0, random_state=1)
>>> lr.fit(X_train_std, y_train)
>>> plot_decision_regions(X_combined_std, y_combined,
...     classifier=lr, test_idx=range(105, 150))
>>> plt.xlabel('petal length [standardized]')
>>> plt.ylabel('petal width [standardized]')
>>> plt.legend(loc='upper left')
>>> plt.show()
```

---

# Logistic Regression

Training a Logistic Regression Model via SciKit-Learn



# Logistic Regression

## Training a Logistic Regression Model via SciKit-Learn

- The probability that training examples belong to a certain class can be computed using the `predict_proba` method. For example, we can predict the probabilities of the first three samples in the test set as follows:

---

```
>>> lr.predict_proba(X_test_std[:3,:])
```

---

- This code snippet returns the following array:

---

```
array([[ 3.81527885e-09,  1.44792866e-01,  8.55207131e-01],  
       [ 8.34020679e-01,  1.65979321e-01,  3.25737138e-13],  
       [ 8.48831425e-01,  1.51168575e-01,  2.62277619e-14]])
```

---

- The first row corresponds to the class-membership probabilities of the first instance, the second row corresponds to the class-membership probabilities of the second instance, and so forth.

# Logistic Regression

## Training a Logistic Regression Model via SciKit-Learn

- Notice that the rows sum all up to one, as expected (you can confirm this by executing `lr.predict_proba(X_test_std[:3, :]).sum(axis=1)`). The highest value in the first row is approximately 0.856, which means that the first sample belongs to class three (Iris-virginica) with a predicted probability of 85.6%.
- So, we can get the predicted class labels by identifying the largest column in each row, for example, using NumPy's `argmax` function:

---

```
>>> lr.predict_proba(X_test_std[:3, :]).argmax(axis=1)
```

---

- The returned class indices are shown here (they correspond to Iris-virginica, Iris-setosa, and Iris-setosa):

---

```
array([2, 0, 0])
```

---



# Logistic Regression

## Training a Logistic Regression Model via SciKit-Learn

- The class labels we obtained from the preceding conditional probabilities is, of course, just a manual approach to calling the `predict` method directly, which we can quickly verify as follows:

---

```
>>> lr.predict(X_test_std[:3, :])  
array([2, 0, 0])
```

---

- Lastly, a word of caution if you want to predict the class label of a single flower sample: scikit-learn expects a two-dimensional array as data input; thus, we have to convert a single row slice into such a format first. One way to convert a single row entry into a two-dimensional data array is to use NumPy's `reshape` method to add a new dimension, as demonstrated here:


---

```
>>> lr.predict(X_test_std[0, :].reshape(1, -1))  
array([2])
```



---

# Overfitting vs Underfitting: Definitions

- **Overfitting** is a common problem in machine learning, where a model performs well on training data but does not generalize well to unseen data (test data).
- **Underfitting** means that our model is not complex enough to capture the pattern in the training data well and therefore also suffers from low performance on unseen data.



# Overfitting vs Underfitting: Examples



# Regularization

- One way of finding a good bias-variance tradeoff is to tune the complexity of the model via regularization.
- Motivation:
  - If a feature is perfectly predictive of the outcome, it will be assigned a very **high weight**.
  - If a feature will attempt to perfectly fit details of the training set could model noisy factors that just accidentally correlate with the class.
- Basic idea:
  - Regularization introduces additional information **to penalize high weights**.
  - Thus a setting of the weights that matches the training data perfectly—but uses many weights with high values to do so—will be penalized more than a setting that matches the data a little less well, but does so using smaller weights
- Regularization is a very useful method to filter out noise from data, and eventually prevent overfitting.

# L2 Regularization

- The most common form of regularization is the so-called *L2* regularization, which can be written as follows:

$$\frac{\lambda}{2} \|\mathbf{w}\|_2^2 = \frac{\lambda}{2} \sum_{j=1}^m w_j^2$$

where  $\lambda$  is the so-called **regularization parameter**.

# Regularization in Logistic Regression

- In order to apply regularization, we just need to add the regularization term to the cost function that we defined for logistic regression to shrink the weights:

$$J(\mathbf{w}) = \left[ \sum_{i=1}^n \left( -y^{(i)} \log(\phi(z^{(i)})) + (1-y^{(i)}) (-\log(1-\phi(z^{(i)}))) \right) \right] + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

- Via the regularization parameter  $\lambda$ , we can then control how well we fit the training data while keeping the weights small. By increasing the value of  $\lambda$ , we increase the regularization strength.
- The parameter  $C$  that is implemented for the Logistic Regression class in scikit-learn is directly related to the regularization parameter  $\lambda$ , which is its inverse:

$$C = \frac{1}{\lambda}$$

# Logistic Regression

## Overfitting vs Underfitting: Regularization

- We can rewrite the regularized cost function of logistic regression as follows:


$$J(\mathbf{w}) = C \left[ \sum_{i=1}^n \left( -\log(\phi(z^{(i)})) + (1 - y^{(i)}) (-\log(1 - \phi(z^{(i)}))) \right) \right] + \frac{1}{2} \|\mathbf{w}\|_2^2$$

- Consequently, decreasing the value of the inverse regularization parameter  $C$  means that we are increasing the regularization strength.

# Logistic Regression

## Overfitting vs Underfitting: Regularization

- We can visualize this consideration by plotting the  $L2$  regularization path for the two weight coefficients:



- As we can see in the resulting plot, the weight coefficients shrink if we decrease the parameter  $C$ , that is, if we increase the regularization strength.



# Logistic Regression

## Overfitting vs Underfitting: Regularization

```
>>> weights, params = [], []
>>> for c in np.arange(-5, 5):
...     lr = LogisticRegression(C=10.**c, random_state=1)
...     lr.fit(X_train_std, y_train)
...     weights.append(lr.coef_[1])
...     params.append(10.**c)
>>> weights = np.array(weights)
>>> plt.plot(params, weights[:, 0], label='petal
length')
>>> plt.plot(params, weights[:, 1], linestyle='--',
label='petal width')
>>> plt.ylabel('weight coefficient')
>>> plt.xlabel('C')
>>> plt.legend(loc='upper left')
>>> plt.xscale('log')
>>> plt.show()
```